Open Ticket AI — Marketing, Sales, Product & Delivery Strategy (Chat Summary)

Executive summary

- **Pricing & economics:** Main offers discussed are **€5k (Pro/Starter)** and **€20k (Enterprise/Rollout)** with **very high gross margin (~99%)**, enabling unusually aggressive spend on lead generation, partner incentives, and “risk reversal” offers.
- **ROI thresholds:** You stated customers reach positive ROI roughly at **~50 tickets/month for €5k** and **~200 tickets/month for €20k**. This becomes the backbone of qualification, messaging, and funnel design.
- **Go-to-market reality:** Classic SEO is currently weak because there is **low explicit search demand** for terms like “Zammad AI / OTOBO AI”. Ranking for head terms (“Zammad”, “OTOBO”) is difficult. Therefore, **direct lead generation (programmatic outbound) + partners** should be the primary acquisition motion, with SEO only as a supporting layer (integration/problem content).
- **Core differentiation:** Since you ship **on-prem** (your USP), customers will have the model bits locally. Therefore, the long-term moat cannot be “weights secrecy”; it must be **system value**: mappings, templates, automations, governance, integrations, and operational reliability.
- **Delivery & monetization:** The simplest commercial gating is **Free vs Premium plugin/template access**, shipped via **Docker images + Docker Compose**, with **private images** for paid tiers. Payment unlock can be handled by a manual approval process issuing a token/licence and registry access.

Key entities and scope discussed

- Company/product: Open Ticket AI
- Ticket system integration focus: Zammad, OTOBO, Znuny
- Delivery/runtime ecosystem: Docker, GitHub (for private registries / distribution)
- LLM integration options: OpenAI (as an optional provider, not required)

1) Pricing, margin, and what it enables

What high margin + €5k/€20k pricing enables

- **Higher CAC tolerance** than typical SaaS: you can afford more spend per qualified opportunity.
- **Partner/referral economics:** paying meaningful referral fees still leaves ample margin.
- **Risk reversal:** pilots, guarantees, credits toward upgrades, and other “low-risk to try” structures.
- **ABM-light:** narrow targeting and higher spend per target account can be rational if the qualification is strong.

ROI-based segmentation (your thresholds)

- **€5k (Pro/Starter)** is compelling at **~50 tickets/month**
- **€20k (Enterprise/Rollout)** is compelling at **~200 tickets/month**

This naturally forms a qualification rule:

- <50 tickets/month: deprioritize (too early)
- 50–199: steer into €5k package
- ≥200: steer to €20k (or €5k as a fast entry if procurement is slow)

2) Go-to-market strategy: why outbound + partners dominate (for now)

Observations you made

- Very few prospects actively search “Zammad AI / OTOBO AI”
- Ranking for “Zammad” / “OTOBO” is highly competitive

Strategic conclusion

- **Primary acquisition:** direct lead generation (programmatic outbound) + partner channel
- **Secondary:** focused SEO that matches *problem intent* and *integration intent* (not head terms)

Channel priority (current state)

Channel	Priority	Why
Programmatic outbound (Python + Google + LLM)	⭐⭐⭐⭐⭐	Creates demand when search demand is low
Partners (integrators/MSPs/consultancies)	⭐⭐⭐⭐⭐	They already sit inside target accounts
Product-led “free download” (open weights)	⭐⭐	High theft/commodity risk in on-prem distribution
SEO (head terms: “Zammad”, “OTOBO”)	⭐	Hard to rank, low direct intent for AI add-ons
SEO (integration/problem content)	⭐⭐⭐	Useful supporting channel; long-term compounding

3) Sales motion and offers (deal flow)

Recommended funnel structure

- **Entry offer:** “48h Ticket Audit” (or similar fast assessment)
- **Conversion offer:** €5k Starter (fast time-to-value)
- **Expansion offer:** €20k Enterprise Rollout

Why this structure fits your market

- Prospects often have the problem but don’t search for it explicitly.
- The audit gives immediate value and reduces perceived risk.
- The €5k offer becomes the “foot-in-the-door” that can expand to €20k.

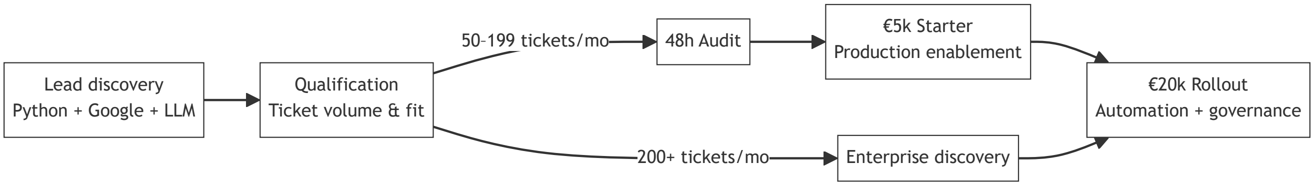

Risk reversal ideas discussed (enabled by margin)

- Credit the €5k into €20k if upgrading within a defined window.
- KPI-based pilot framing (clear success criteria).
- Clear “what you get in X days” packaging.

4) Resource allocation: development-first, sales handled by partner

Your constraint

- You want **more time in development**
- **Constanze handles sales**
- Need to decide whether to invest money into freelancers/agency

Conclusion reached

- Avoid generic marketing agencies (especially pure SEO/content agencies) at this stage.
- Put budget into **assets that scale**:
  - improving the **lead discovery script**
  - enrichment + scoring + outreach automation tooling
  - partner enablement materials (templates, docs, bundles)

Practical spend idea

- Hire a **strong Python/data automation freelancer** to turn your lead script into a repeatable “Lead OS”.
- Optionally later add an outbound SDR/assistant for list cleaning + first-touch, but your script is the main leverage.

5) Product architecture (runtime, plugins, templates, studio)

Components you described

- **Tagging models by tier**
  - Free: Qwen3 0.6B, ~100 tags
  - Pro: Qwen3 4B, ~200 tags
  - Enterprise: Qwen3 14B, ~400 tags (still under consideration for broader suite)
- **Python runtime program**
  - Loads the model
  - Applies thresholding logic (e.g., family highest confidence)
- **OTAI Runtime Core (LGPL)**

  - Flow/pipeline/workflow engine
- **OTAI Tag Mapping plugin**
  - Maps tags to downstream values via CEL-defined mapping rules
- **Zammad integration plugin**
  - Integration exists, currently beta
- **Template plugins**
  - Provide prebuilt flows for: queue/group assignment, priority, SLA, custom fields, etc.
- **OTAI Studio / Console / Control Plane**
  - Mapping Studio generates mapping YAML configs from queues/priorities/etc.
  - Deployment model decided: **Enterprise local**, **Free/Pro SaaS**

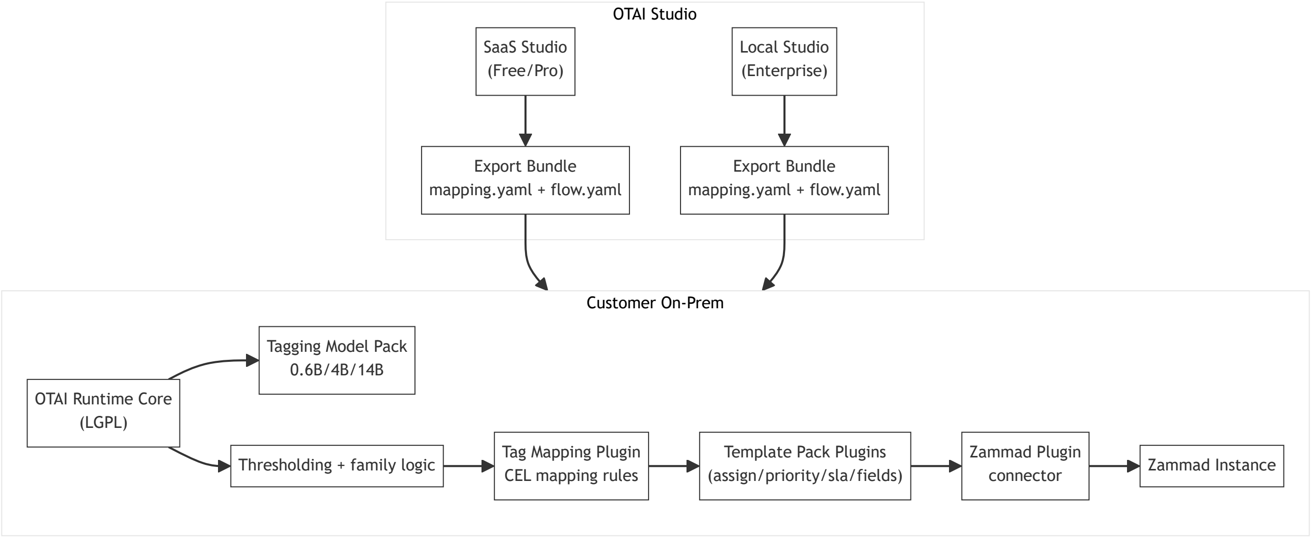
OTAI Studio
SaaS Studio (Free/Pro)
Local Studio (Enterprise)
Export Bundle mapping.yaml + flow.yaml
Export Bundle mapping.yaml + flow.yaml
Customer On-Prem
OTAI Runtime Core (LGPL)
Tagging Model Pack 0.6B/4B/14B
Thresholding + family logic
Tag Mapping Plugin CEL mapping rules
Template Pack Plugins (assign/priority/sla/fields)
Zammad Plugin connector
Zammad Instance

6) Free vs Pro vs Enterprise strategy (and “theft” concern)

Core issue you raised

- On-prem means the **model can be copied** once delivered.
- A naive “free model download” can accelerate competitors.

Important implication

- Because you ship on-prem, the moat must be **beyond weights**:
  - templates, mapping logic UX, integration quality, governance, deployment simplicity, updates

Simplest gating strategy agreed

- **Gate by plugin/template availability**, not by ticket counts.
- Implement Free vs Premium as separate plugin packages (Free plugin vs Premium plugin).

Suggested tiering (high-level)

Tier	Models/Tags	Templates/Actions	Studio
Free	0.6B / ~100 tags	only basic queue/group template	SaaS Studio
Pro (€5k)	0.6B/4B, more tags, more quantizations	+ priority + custom fields + richer routing templates	SaaS Studio
Enterprise (€20k)	up to 14B / ~400 tags + suite	+ SLA + governance + audits + multi-flow	Local Studio

7) Delivery & access control (Docker/Compose + plugins)

Your preferred delivery

- **Docker images + Docker Compose** for simplest customer install.

Main design decision

- Don’t load premium Python packages at runtime via pip if you can avoid it.
- Prefer **private Docker images** for premium packs (and/or runtime-pro images).

Recommended distribution pattern

- **Public images**: runtime + free plugins
- **Private images**: premium plugin packs, premium model packs, enterprise extras
- Access control:
  - Customer gets registry access after payment
  - Customer gets a licence/token (optional but recommended)
  - Compose pulls the private images

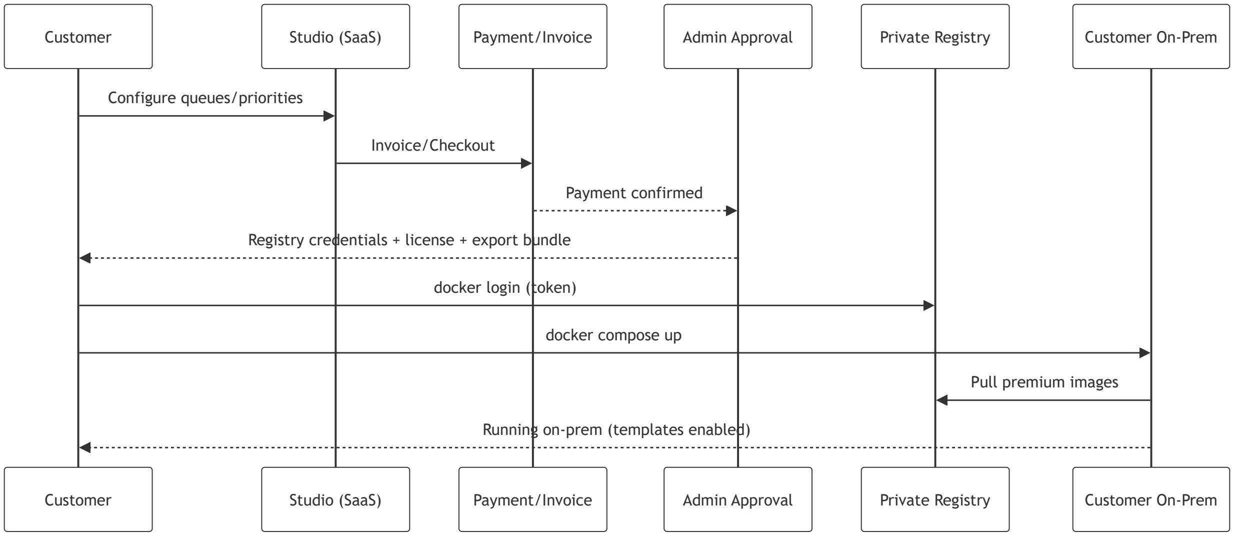

Why “API key loads packages” is weaker

- More operational complexity (pip, proxies, TLS, retries, caching)
- More failure modes for customers
- Harder to support air-gapped scenarios

8) Mapping Studio: SaaS vs On-Prem and “On-Prem USP confusion”

Concern you raised

- On-prem is your USP; SaaS Studio might confuse or scare buyers.

Resolution

- Message it as: **runtime + inference always on-prem**; Studio is just a **YAML generator**.
- Offer both modes:
  - **SaaS for Free/Pro** (fast onboarding)
  - **Local Studio for Enterprise** (compliance expectation)

Practical UX mitigation

- Provide “No storage mode” in SaaS (generate → export; minimal persistence).
- Make the export bundle identical between SaaS and local Studio so customers see it as “the same product”.

9) Mapping generation with LLM (SaaS and On-Prem)

Your plan

- Use an LLM to generate mapping YAML/CEL based on:
  - queues + descriptions
  - priorities + descriptions
  - desired assignments/automation logic

On-Prem mapping generation options discussed

- Use a local model (e.g., 14B) for higher quality mapping drafts.
- Alternatively allow a “provider” approach:
  - local model
  - customer-provided key (OpenAI/Azure) if allowed
  - template/rule-based fallback

Key technical insight

- Mapping generation quality depends heavily on:
  - strong schema constraints
  - allowed-value lists (queues/priorities/fields)
  - validation + review UI

    - template scaffolding (avoid “free-form CEL hallucinations”)

10) Open questions and decisions pending

Decisions pending

- Whether Enterprise is “just 14B + 400 tags” or a broader automation/governance suite.
- How strict SaaS Studio persistence should be for Pro customers (persist vs no-storage mode).
- Whether you want any public “free model” distribution beyond the simplest baseline.

Potential future improvements

- Enterprise-grade governance:
    - mapping versioning, approvals, audit reports
- Partner enablement:
    - standardized templates per ticket system
- Lead OS:
    - scoring model for “ticket volume likelihood”
    - automated enrichment and personalized outreach drafts

Recommended “next concrete steps”

1. **Define SKU boundaries by templates/actions**
    - Free: only queue/group assignment template
    - Pro: + priority + custom fields + richer routing templates
    - Enterprise: + SLA + governance + multi-flow packs
2. **Standardize the Studio export bundle**
    - same files for SaaS and local Studio
3. **Implement premium distribution via private images**
    - keep Python plugins inside images
4. **Strengthen the mapping generator**
    - schema constraints + validator + “draft → review → export”
5. **Invest budget into the lead engine**
    - scale the Python + Google + LLM pipeline
    - add scoring + segmentation using your ROI thresholds